# Edimah SYNESIUS SONGO

## Mathematical statistician – AI alignment & evaluation

linkedin.com/in/edimah-synesius github.com/edimah
edimah.synesius-songo@proton.me
edimah.github.io
Paris


## </> Projects

### Predicting false collusion alerts under incomplete information — October 2026

AI Collusion Research Sprint, Apart Research, solo

- Simulated a market where a standard screen's false-accusation rate for collusion rises 5.8% → 100% as auditor information degrades.
- Built a simple analytical model that predicts these rates closely, without being fitted to them.
- Proposed a "re-asking" test that re-queries each agent on its recorded inputs, without the secret channel, to tell hidden collusion apart from simply shared responses.
- Ran an exploratory pilot for LLM agents: 0/9 false flags, 3/3 planted pairs detected (3 games; statistic chosen after the pre-registered one failed).

Apart Research project page · github.com/Edimah/hackathon-ai-collusion

Hypothesis Testing | Sensitivity Analysis | Mathematical Modelling | Simulation | Python

### Detecting agentic collusion in public logs with statistical methods — September 2026

AI Incident Response Sprint, Apart Research × CeSIA, Track 2, solo

- Used statistical methods to detect an LLM-agent swarm from public edit logs alone.
- Reconstructed the logs from two public sources over 2003–2026.
- Designed a ladder of four null models to characterise the incident and its statistical significance.
- Built a weekly check as the incident-response deliverable.
- It flags the three largest incident weeks and none of the 2,236 earlier wiki-weeks (2016–2025).

Apart Research project page · github.com/Edimah/agentic-collusion-detector-in-public-logs

Change-Point Detection | Hypothesis Testing | Statistical Modelling | Python

### A Calibrated Two-Level Detector for Secret Loyalties – Application to an LLM — July 2026

"Secret Loyalties" hackathon, Apart Research × Formation Research, team of 2

- Designed an evaluation framework, then measured its false-positive rate on a clean model.
- Using matched-control calibration, cut false positives from 12/16 to 2/16.
- Result: top 25% of 179 submissions; one reviewer cited the calibration as "a key contribution".

Apart Research project page · github.com/edimah/secret-loyalties
Write-up: edimah.github.io/blog/2026/secret-loyalties-outputs-only/

Hypothesis Testing | Statistical Calibration | LLM Evaluation

### Advanced methods for model calibration with uncertain inputs — April – September 2023

M2 research thesis, EDF R&D, Chatou, solo

- Topic: Bayesian calibration of physical simulation models against noisy experimental data.
- Designed adaptive MCMC-based re-adjustment strategies.
- Designed calibration updates that respect the metrologists' measurement chain.

Bayesian Calibration | MCMC | Uncertainty Quantification | Python | OpenTURNS

## Experience

**Present – October 2025** — **Independent Consultant, Applied Statistics, AI Implementation**

- Statistics & Data consulting practice
- Flagship engagement: interim **Data Lead** at **CRCDC Guadeloupe** since February 2026.
- Rebuilding the regional cancer-screening data chain to deliver reliable data to stakeholders.
- Delivered a data-chain audit and methodology documentation under regulatory constraints.
- Assessed the feasibility of AI tools for pharmacy workflows (regulated health care).

Data Management | Communication | R | SQL | eSIS | DataQuery | Data Visualisation | Health Data Regulation

**September 2024 – November 2023** — **Doctoral Researcher, IFP Energies Nouvelles / ONERA, Île-de-France**

- Multidisciplinary sensitivity analysis for complex physics systems.
- Development of sensitivity indices and explainability metrics in R and Python.
- Left the doctorate before completion.

R | Python | Sensitivity Analysis | Uncertainty Quantification

September 2023 – April 2023 | **Research Intern, EDF R&D, Chatou**
› **M2 Thesis**: *Advanced methods for model calibration with uncertain inputs* (see Projects).
Python · OpenTURNS · Bayesian Inference · Computational Statistics · MCMC

August 2022 – July 2022 | **Data Analyst Intern, Hôpital Avicenne, Bobigny**
› **Data analysis** of patient attendance and trend identification.
Tableau · Data Analysis · Data Visualisation · R · SQL

## Education

2023 **Master 2 – MASH** (Mathematics, Machine Learning, Sciences, Humanities) at **Université PSL** (Dauphine & ENS), with Merit (mention Bien).
Research thesis at EDF R&D: *Advanced methods for model calibration with uncertain inputs*.
2022 **Master 1 – Applied Mathematics**, Statistics track, at **Paris Dauphine-PSL**.
2021 **Bachelor's – MEFA** (Mathematics, Economics, Finance, Actuarial Science) at **Paris Dauphine-PSL**
*2025* *TOEFL iBT (112/120)*

## AI Safety Engagement

**Courses** ARENA material, self-study, started 14 Oct 2026 — github.com/Edimah/arena
Lens Academy — **"AI Risk Fundamentals"** intensive, certified Oct 2026
BlueDot Impact — **"Future of AI"** course, certified Oct 2026
BlueDot Impact — AGI Strategy, self-paced, from Oct 2026

**Reflections** A couple of AI safety posts on my personal blog, edimah.github.io; one crossposted to the Effective Altruism Forum

**Sprints & hackathons** Three Apart Research sprints (July–October 2026, cf. "Projects")

**Community** Slack: EA France, AI Alignment · Discord: AI Safety & SLT, Sûreté de l'IA, Apart, Lens Academy
Paris AI Safety Community ("Nouvelle Machine"): dinners and workshops, participant since Oct 2026 · Effective Altruism France events; registered for EA Summit Paris, 5 Dec 2026

## Skills

**Data** Data management · Data visualisation · Health data regulation · Data Analysis

**Applied stats** Statistical modelling · Hypothesis testing · Bayesian inference
Computational statistics, MCMC · Monte Carlo simulation

**Applied maths** Mathematical modelling · Optimisation · Uncertainty quantification
Sensitivity analysis · Bayesian calibration

**Programming** Python · R · SQL · C++ · Git, Make, LaTeX

**Languages** French (native) · English (C2, TOEFL 112/120) · Spanish (B2)